

SUMMARY

Pre-final year ECE student specializing in Applied AI and ML Engineering, with hands-on experience building RAG pipelines, LLM-powered applications, and full-stack AI systems. Skilled in Python, FastAPI, LangChain, and vector databases, with a focus on production-ready AI architectures

WORK EXPERIENCE

Data Engineering Intern | SPRINGER CAPITAL | January 2026 - April 2026

- Built and orchestrated 2+ end-to-end ETL pipelines using Python, Apache Airflow, and PostgreSQL for structured data ingestion and processing
- Designed data transformation workflows converting raw CSV data into analytics-ready Parquet and SQL datasets (Bronze → Silver architecture)

Data Analytics Intern | INFOSYS SPRINGBOARD | [LINK](#) October 2025 - December 2025

- Analyzed 6.7M+ flight records using Python and SQL to identify delay patterns across weather, carrier operations, and seasonal trends.
- Built analytical queries and feature engineering pipelines to investigate data issues and generate insights, identifying 7.4× higher winter cancellations and 2.8× holiday delay impact.

PROJECTS

CreatorLens (Next.js, Langchain, FastAPI) | [LINK](#) December 2024 - February 2025

- Built a multimodal RAG pipeline ingesting YouTube/Instagram transcripts, metadata, and engagement stats into a vector store, enabling semantic retrieval across 2 video sources with LLM-powered comparative analysis via a Next.js chat interface
- Architected a 3-tier transcript extraction pipeline (yt-dlp → authenticated youtube-transcript-api → unauthenticated fallback) achieving ~95% ingestion success rate against rate limiting and bot detection
- Designed modular FastAPI backend with RAG orchestration (indexer, retriever, prompt chain) and Pydantic schema validation, reducing invalid API responses to 0 across all ingestion workflows

AutoFlow– Visual Workflow Automation Platform (Next.js, FastAPI)| [LINK](#) August 2025 – October 2025

- Designed NL→ workflow generation pipeline using Groq LLM with structured JSON output, Pydantic schema validation, and an auto-retry loop for malformed responses
- Built a node execution engine that pipes data sequentially between AI, communication, and webhook nodes with real-time monitoring
- Built a responsive frontend using Next.js and React for an AI workflow automation platform, enabling users to design and execute automated pipelines through an intuitive interface

LLM Evaluation Platform (Python, Groq, GitHub Actions, Docker) | [LINK](#) December 2024 - February 2025

- Built an automated LLM evaluation pipeline comparing prompt versions against a 100+ case golden dataset, detecting accuracy regressions across configurable warning/critical drift thresholds with category-level breakdowns
- Implemented async classification engine using asyncio Semaphore for concurrent Groq API calls, cutting evaluation runtime from 15+ minutes to under 3 minutes, with CI/CD via GitHub Actions, Dockerized deployment, and Slack regression alerts

SKILLS SUMMARY

- AI & ML :** RAG Pipelines, LLMs, Prompt Engineering, Vector Databases (FAISS/Chroma), Embeddings, Groq, LLM Evaluation, XGBoost, Scikit-learn, Pydantic, Langchain
- Data Engineering:** ETL Pipelines, Apache Airflow, PostgreSQL, Data Transformation, Parquet
- Web Development & Tools:** Next JS, Fast API, MongoDB, Git, GitHub, JIRA, Docker, Power BI, Tableau, DAX, VS Code
- Programming:** Python, JavaScript, SQL, C

EDUCATION

Christ College Engineering

B.E. Electronics and Communication Engineering ; GPA: 8.06
Machine Learning (MINOR) , English Proficiency: C1 (Advanced)

June 2023 - August 2027
Kerala, India

- Achieved Top 2 rank in class (2 semesters)

CERTIFICATES & ACHIEVEMENTS

Generative AI Engineering (IBM) | [CERTIFICATE](#)

December 2025

2× Hackathon Winner: Qbit (Jain University)

HackQuest (Rajagiri college of engineering)